

```
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
import os
import pandas as pd
import seaborn as sns
import random
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten,
Dense, Dropout, Activation, BatchNormalization
from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping
from tensorflow.keras.preprocessing.image import ImageDataGenerator,
load_img, img_to_array
from tensorflow.keras.optimizers import Adam
from keras.utils import plot_model
from PIL import Image
from collections import Counter
from glob import glob
from tensorflow.keras.utils import plot_model
```

Loading The Data

```
train_path = 'fruits-360_dataset/fruits-360/Training/'
test_path = 'fruits-360_dataset/fruits-360/Test/'
```

Dataset Overview

Title: Fruit and Vegetable Image Recognition

Source: [Kaggle Fruits- 360 Dataset](#)

Dataset Characteristics:

- Total number of images: 90,483
- Number of classes: 131
- Image dimensions: 100x100 pixels
- Training set size: 67,692 images
- Test set size: 22,688 images

Credit for base CNN: <https://medium.com/hackerdawn/fruit-image-classification-using-cnn-on-google-colab-4fe7274418a5>

```

img = load_img(train_path + "Apple Braeburn/0_100.jpg",
target_size=(100,100))
plt.imshow(img)
plt.axis("off")
plt.show()

x = img_to_array(img)
print(x.shape)

```



```

(100, 100, 3)

images = ['Apple Crimson Snow', 'Banana', 'Apricot', 'Blueberry',
'Corn', 'Kiwi', 'Pear', 'Watermelon', 'Orange']
fig = plt.figure(figsize =(10,5))
for i in range(len(images)):
    ax = fig.add_subplot(3,3,i+1,xticks=[],yticks=[])
    plt.title(images[i])
    plt.axis("off")
    ax.imshow(load_img(train_path + images[i] + "/0_100.jpg",
target_size=(100,100)))

```

Apple Crimson Snow



Blueberry



Pear



Banana



Corn



Watermelon



Apricot



Kiwi



Orange



```
fruits = []
fruits_image = []
for i in os.listdir(train_path):
    for image_filename in os.listdir(train_path + i):
        fruits.append(i)
        fruits_image.append(i + '/' + image_filename)
```

```
newData = Counter(fruits)
frequent_fruits = newData.most_common(20)
print("Top 20 frequent Fruits:")
frequent_fruits
```


Top 20 frequent Fruits:

```
[('Grape Blue', 984),
 ('Plum 3', 900),
 ('Cherry 2', 738),
 ('Cherry Rainier', 738),
 ('Melon Piel de Sapo', 738),
 ('Peach 2', 738),
 ('Strawberry Wedge', 738),
 ('Tomato 1', 738),
 ('Tomato 3', 738),
 ('Walnut', 735),
 ('Pear Stone', 711),
 ('Cauliflower', 702),
```

```
( 'Fig', 702),
( 'Pear Forelle', 702),
( 'Pepper Orange', 702),
( 'Pear 2', 696),
( 'Tomato Heart', 684),
( 'Apple Red Yellow 2', 672),
( 'Tomato 2', 672),
( 'Pear Red', 666)]
```

```
fruit_names = [fruit[0] for fruit in frequent_fruits]
fruit_counts = [fruit[1] for fruit in frequent_fruits]
```

```
plt.figure(figsize=(12, 6))
plt.bar(fruit_names, fruit_counts, color='skyblue')
plt.xlabel('Fruit')
plt.ylabel('Frequency')
plt.title('Top 20 Most Frequent Fruits in the Dataset')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



Model

```
className = glob(train_path + '/*')
number_of_class = len(className)
```

```

model = Sequential()
model.add(Conv2D(32,(3,3),input_shape = x.shape))
model.add(Activation("relu"))
model.add(MaxPooling2D())

model.add(Conv2D(32,(3,3)))
model.add(Activation("relu"))
model.add(MaxPooling2D())

model.add(Conv2D(64,(3,3)))
model.add(Activation("relu"))
model.add(MaxPooling2D())

model.add(Flatten())
model.add(Dense(1024))
model.add(Activation("relu"))
model.add(Dropout(0.5))
model.add(Dense(number_of_class))
model.add(Activation("softmax"))

model.compile(loss = "categorical_crossentropy",
optimizer = "rmsprop",
metrics = ["accuracy"])
model.summary()

```

Model: "sequential_1"

Layer (type) Param #	Output Shape	
conv2d_3 (Conv2D) 896	(None, 98, 98, 32)	
activation_5 (Activation) 0	(None, 98, 98, 32)	
max_pooling2d_3 (MaxPooling2D) 0	(None, 49, 49, 32)	
conv2d_4 (Conv2D) 9,248	(None, 47, 47, 32)	
activation_6 (Activation)	(None, 47, 47, 32)	

0				
		max_pooling2d_4 (MaxPooling2D)	(None, 23, 23, 32)	
0				
		conv2d_5 (Conv2D)	(None, 21, 21, 64)	
18,496				
		activation_7 (Activation)	(None, 21, 21, 64)	
0				
		max_pooling2d_5 (MaxPooling2D)	(None, 10, 10, 64)	
0				
		flatten_1 (Flatten)	(None, 6400)	
0				
		dense_2 (Dense)	(None, 1024)	
6,554,624				
		activation_8 (Activation)	(None, 1024)	
0				
		dropout_1 (Dropout)	(None, 1024)	
0				
		dense_3 (Dense)	(None, 131)	
134,275				
		activation_9 (Activation)	(None, 131)	
0				

Total params: 6,717,539 (25.63 MB)

Trainable params: 6,717,539 (25.63 MB)

Non-trainable params: 0 (0.00 B)

```
epochs = 100
batch_size = 64
```

Augmented Data

```
train_datagen = ImageDataGenerator(
    rescale=1./255,
    shear_range=0.3,
    horizontal_flip=True,
    zoom_range=0.3
)

test_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(
    directory = train_path,
    target_size= x.shape[:2],
    batch_size = batch_size,
    color_mode= "rgb",
    class_mode= "categorical"
)

test_generator = test_datagen.flow_from_directory(
    directory = test_path,
    target_size= x.shape[:2],
    batch_size = batch_size,
    color_mode= "rgb",
    class_mode= "categorical"
)

Found 67692 images belonging to 131 classes.
Found 22688 images belonging to 131 classes.
```

Fitting the model

```
hist = model.fit(
    x=train_generator,
    steps_per_epoch=1600 // batch_size,
    epochs=epochs,
    validation_data=test_generator,
    validation_steps=800 // batch_size
)

Epoch 1/100
```

```
c:\Users\adamn\AppData\Local\Programs\Python\Python311\Lib\site-  
packages\keras\src\trainers\data_adapters\py_dataset_adapter.py:120:  
UserWarning: Your `PyDataset` class should call  
`super().__init__(**kwargs)` in its constructor. `**kwargs` can  
include `workers`, `use_multiprocessing`, `max_queue_size`. Do not  
pass these arguments to `fit()`, as they will be ignored.  
    self._warn_if_super_not_called()
```

```
25/25 ━━━━━━━━━━━ 31s 1s/step - accuracy: 0.0155 - loss:  
4.9032 - val_accuracy: 0.0339 - val_loss: 4.6321  
Epoch 2/100  
25/25 ━━━━━━━━━━━ 22s 896ms/step - accuracy: 0.0357 - loss:  
4.5676 - val_accuracy: 0.0911 - val_loss: 3.8680  
Epoch 3/100  
25/25 ━━━━━━━━━━━ 20s 816ms/step - accuracy: 0.0756 - loss:  
3.9581 - val_accuracy: 0.1536 - val_loss: 3.3425  
Epoch 4/100  
25/25 ━━━━━━━━━━━ 16s 645ms/step - accuracy: 0.1032 - loss:  
3.6045 - val_accuracy: 0.2148 - val_loss: 2.9657  
Epoch 5/100  
25/25 ━━━━━━━━━━━ 16s 660ms/step - accuracy: 0.1923 - loss:  
3.1241 - val_accuracy: 0.2917 - val_loss: 2.6141  
Epoch 6/100  
25/25 ━━━━━━━━━━━ 19s 787ms/step - accuracy: 0.2547 - loss:  
2.7536 - val_accuracy: 0.4583 - val_loss: 2.1070  
Epoch 7/100  
25/25 ━━━━━━━━━━━ 17s 668ms/step - accuracy: 0.2606 - loss:  
2.6075 - val_accuracy: 0.5560 - val_loss: 1.7412  
Epoch 8/100  
25/25 ━━━━━━━━━━━ 16s 663ms/step - accuracy: 0.3575 - loss:  
2.1925 - val_accuracy: 0.5664 - val_loss: 1.6075  
Epoch 9/100  
25/25 ━━━━━━━━━━━ 17s 680ms/step - accuracy: 0.3864 - loss:  
2.0316 - val_accuracy: 0.5365 - val_loss: 1.6849  
Epoch 10/100  
25/25 ━━━━━━━━━━━ 16s 647ms/step - accuracy: 0.4462 - loss:  
1.8306 - val_accuracy: 0.6432 - val_loss: 1.3559  
Epoch 11/100  
25/25 ━━━━━━━━━━━ 16s 654ms/step - accuracy: 0.5071 - loss:  
1.5842 - val_accuracy: 0.7266 - val_loss: 1.0184  
Epoch 12/100  
25/25 ━━━━━━━━━━━ 16s 637ms/step - accuracy: 0.5518 - loss:  
1.4552 - val_accuracy: 0.6784 - val_loss: 1.0842  
Epoch 13/100  
25/25 ━━━━━━━━━━━ 16s 645ms/step - accuracy: 0.6137 - loss:  
1.2439 - val_accuracy: 0.7604 - val_loss: 0.7982  
Epoch 14/100  
25/25 ━━━━━━━━━━━ 15s 622ms/step - accuracy: 0.6011 - loss:  
1.2616 - val_accuracy: 0.6771 - val_loss: 1.0708  
Epoch 15/100
```


25/25 _____ 15s 611ms/step - accuracy: 0.6136 - loss: 1.2278 - val_accuracy: 0.7839 - val_loss: 0.7646
Epoch 16/100
25/25 _____ 16s 651ms/step - accuracy: 0.6506 - loss: 1.0628 - val_accuracy: 0.8346 - val_loss: 0.6399
Epoch 17/100
25/25 _____ 17s 698ms/step - accuracy: 0.6981 - loss: 0.9152 - val_accuracy: 0.8255 - val_loss: 0.6009
Epoch 18/100
25/25 _____ 17s 710ms/step - accuracy: 0.7365 - loss: 0.8702 - val_accuracy: 0.8385 - val_loss: 0.5897
Epoch 19/100
25/25 _____ 17s 686ms/step - accuracy: 0.7089 - loss: 0.8786 - val_accuracy: 0.7500 - val_loss: 0.7517
Epoch 20/100
25/25 _____ 17s 681ms/step - accuracy: 0.6862 - loss: 0.9383 - val_accuracy: 0.8333 - val_loss: 0.5603
Epoch 21/100
25/25 _____ 16s 653ms/step - accuracy: 0.7594 - loss: 0.7431 - val_accuracy: 0.8516 - val_loss: 0.4850
Epoch 22/100
25/25 _____ 14s 568ms/step - accuracy: 0.7745 - loss: 0.7082 - val_accuracy: 0.7357 - val_loss: 0.8744
Epoch 23/100
25/25 _____ 15s 598ms/step - accuracy: 0.7620 - loss: 0.7071 - val_accuracy: 0.8685 - val_loss: 0.4581
Epoch 24/100
25/25 _____ 14s 588ms/step - accuracy: 0.8019 - loss: 0.6223 - val_accuracy: 0.8568 - val_loss: 0.4625
Epoch 25/100
25/25 _____ 14s 565ms/step - accuracy: 0.8242 - loss: 0.5616 - val_accuracy: 0.8542 - val_loss: 0.4282
Epoch 26/100
25/25 _____ 14s 553ms/step - accuracy: 0.7786 - loss: 0.6924 - val_accuracy: 0.8828 - val_loss: 0.3729
Epoch 27/100
25/25 _____ 14s 558ms/step - accuracy: 0.8177 - loss: 0.5418 - val_accuracy: 0.8789 - val_loss: 0.3964
Epoch 28/100
25/25 _____ 13s 531ms/step - accuracy: 0.8086 - loss: 0.6033 - val_accuracy: 0.7956 - val_loss: 0.6671
Epoch 29/100
25/25 _____ 13s 542ms/step - accuracy: 0.8144 - loss: 0.5980 - val_accuracy: 0.8477 - val_loss: 0.5070
Epoch 30/100
25/25 _____ 13s 535ms/step - accuracy: 0.8289 - loss: 0.5321 - val_accuracy: 0.8534 - val_loss: 0.4676
Epoch 31/100

```
c:\Users\adamn\AppData\Local\Programs\Python\Python311\Lib\
contextlib.py:155: UserWarning: Your input ran out of data;
interrupting training. Make sure that your dataset or generator can
generate at least `steps_per_epoch * epochs` batches. You may need to
use the `.repeat()` function when building your dataset.
    self.gen.throw(typ, value, traceback)
```

```
25/25 _____ 16s 634ms/step - accuracy: 0.8224 - loss:
0.5218 - val_accuracy: 0.8854 - val_loss: 0.3928
Epoch 32/100
```

```
25/25 _____ 14s 559ms/step - accuracy: 0.8543 - loss:
0.4470 - val_accuracy: 0.9115 - val_loss: 0.3055
Epoch 33/100
```

```
25/25 _____ 13s 524ms/step - accuracy: 0.8776 - loss:
0.3978 - val_accuracy: 0.8958 - val_loss: 0.3904
Epoch 34/100
```

```
25/25 _____ 13s 543ms/step - accuracy: 0.8724 - loss:
0.4136 - val_accuracy: 0.9036 - val_loss: 0.3497
Epoch 35/100
```

```
25/25 _____ 15s 615ms/step - accuracy: 0.8716 - loss:
0.3837 - val_accuracy: 0.9271 - val_loss: 0.2213
Epoch 36/100
```

```
25/25 _____ 15s 627ms/step - accuracy: 0.8747 - loss:
0.3846 - val_accuracy: 0.9336 - val_loss: 0.2214
Epoch 37/100
```

```
25/25 _____ 15s 608ms/step - accuracy: 0.8872 - loss:
0.3367 - val_accuracy: 0.8763 - val_loss: 0.3846
Epoch 38/100
```

```
25/25 _____ 14s 580ms/step - accuracy: 0.8426 - loss:
0.4743 - val_accuracy: 0.9089 - val_loss: 0.3090
Epoch 39/100
```

```
25/25 _____ 15s 630ms/step - accuracy: 0.8862 - loss:
0.3530 - val_accuracy: 0.9089 - val_loss: 0.2730
Epoch 40/100
```

```
25/25 _____ 16s 644ms/step - accuracy: 0.8989 - loss:
0.3220 - val_accuracy: 0.8945 - val_loss: 0.3653
Epoch 41/100
```

```
25/25 _____ 15s 618ms/step - accuracy: 0.8721 - loss:
0.4229 - val_accuracy: 0.9049 - val_loss: 0.2685
Epoch 42/100
```

```
25/25 _____ 536s 22s/step - accuracy: 0.8933 - loss:
0.3362 - val_accuracy: 0.9401 - val_loss: 0.2523
Epoch 43/100
```

```
25/25 _____ 6s 198ms/step - accuracy: 0.8661 - loss:
0.4096 - val_accuracy: 0.8828 - val_loss: 0.3680
Epoch 44/100
```

```
25/25 _____ 25s 858ms/step - accuracy: 0.8987 - loss:
0.3102 - val_accuracy: 0.9323 - val_loss: 0.2154
Epoch 45/100
```

```
25/25 _____ 21s 869ms/step - accuracy: 0.9102 - loss:
```

0.2597 - val_accuracy: 0.9167 - val_loss: 0.3112
Epoch 46/100
25/25 _____ 30s 1s/step - accuracy: 0.9079 - loss:
0.2747 - val_accuracy: 0.9297 - val_loss: 0.2825
Epoch 47/100
25/25 _____ 23s 914ms/step - accuracy: 0.8993 - loss:
0.3322 - val_accuracy: 0.9167 - val_loss: 0.3388
Epoch 48/100
25/25 _____ 19s 771ms/step - accuracy: 0.9003 - loss:
0.2969 - val_accuracy: 0.9049 - val_loss: 0.2696
Epoch 49/100
25/25 _____ 19s 766ms/step - accuracy: 0.9134 - loss:
0.2714 - val_accuracy: 0.8034 - val_loss: 0.6712
Epoch 50/100
25/25 _____ 20s 791ms/step - accuracy: 0.9180 - loss:
0.2679 - val_accuracy: 0.9245 - val_loss: 0.2539
Epoch 51/100
25/25 _____ 21s 846ms/step - accuracy: 0.8887 - loss:
0.3421 - val_accuracy: 0.9062 - val_loss: 0.3192
Epoch 52/100
25/25 _____ 18s 712ms/step - accuracy: 0.9032 - loss:
0.2909 - val_accuracy: 0.9154 - val_loss: 0.2352
Epoch 53/100
25/25 _____ 24s 985ms/step - accuracy: 0.9031 - loss:
0.2764 - val_accuracy: 0.8112 - val_loss: 0.6436
Epoch 54/100
25/25 _____ 24s 981ms/step - accuracy: 0.9111 - loss:
0.2893 - val_accuracy: 0.9362 - val_loss: 0.2041
Epoch 55/100
25/25 _____ 20s 811ms/step - accuracy: 0.9328 - loss:
0.2052 - val_accuracy: 0.9453 - val_loss: 0.1981
Epoch 56/100
25/25 _____ 19s 764ms/step - accuracy: 0.8929 - loss:
0.2724 - val_accuracy: 0.8958 - val_loss: 0.3159
Epoch 57/100
25/25 _____ 18s 708ms/step - accuracy: 0.9323 - loss:
0.2121 - val_accuracy: 0.9336 - val_loss: 0.2355
Epoch 58/100
25/25 _____ 17s 675ms/step - accuracy: 0.9186 - loss:
0.2236 - val_accuracy: 0.9349 - val_loss: 0.1716
Epoch 59/100
25/25 _____ 16s 635ms/step - accuracy: 0.9421 - loss:
0.1598 - val_accuracy: 0.9505 - val_loss: 0.2241
Epoch 60/100
25/25 _____ 16s 639ms/step - accuracy: 0.9295 - loss:
0.2465 - val_accuracy: 0.9639 - val_loss: 0.1678
Epoch 61/100
25/25 _____ 20s 829ms/step - accuracy: 0.9204 - loss:
0.2243 - val_accuracy: 0.9414 - val_loss: 0.1853

Epoch 62/100
25/25 _____ 18s 751ms/step - accuracy: 0.9311 - loss: 0.2061 - val_accuracy: 0.9557 - val_loss: 0.1570
Epoch 63/100
25/25 _____ 16s 663ms/step - accuracy: 0.9280 - loss: 0.2147 - val_accuracy: 0.9544 - val_loss: 0.1568
Epoch 64/100
25/25 _____ 14s 595ms/step - accuracy: 0.9178 - loss: 0.2469 - val_accuracy: 0.9492 - val_loss: 0.1677
Epoch 65/100
25/25 _____ 14s 593ms/step - accuracy: 0.9284 - loss: 0.2176 - val_accuracy: 0.9271 - val_loss: 0.2306
Epoch 66/100
25/25 _____ 15s 615ms/step - accuracy: 0.9146 - loss: 0.2355 - val_accuracy: 0.9583 - val_loss: 0.1645
Epoch 67/100
25/25 _____ 15s 592ms/step - accuracy: 0.9415 - loss: 0.1631 - val_accuracy: 0.9453 - val_loss: 0.1972
Epoch 68/100
25/25 _____ 15s 598ms/step - accuracy: 0.9384 - loss: 0.1988 - val_accuracy: 0.9727 - val_loss: 0.1101
Epoch 69/100
25/25 _____ 15s 606ms/step - accuracy: 0.9414 - loss: 0.1780 - val_accuracy: 0.9414 - val_loss: 0.2507
Epoch 70/100
25/25 _____ 15s 601ms/step - accuracy: 0.9380 - loss: 0.1702 - val_accuracy: 0.9661 - val_loss: 0.0969
Epoch 71/100
25/25 _____ 15s 597ms/step - accuracy: 0.9414 - loss: 0.1788 - val_accuracy: 0.9453 - val_loss: 0.2345
Epoch 72/100
25/25 _____ 14s 571ms/step - accuracy: 0.9382 - loss: 0.2138 - val_accuracy: 0.9440 - val_loss: 0.1587
Epoch 73/100
25/25 _____ 18s 713ms/step - accuracy: 0.9554 - loss: 0.1440 - val_accuracy: 0.9596 - val_loss: 0.1454
Epoch 74/100
25/25 _____ 15s 610ms/step - accuracy: 0.9384 - loss: 0.1797 - val_accuracy: 0.9583 - val_loss: 0.1718
Epoch 75/100
25/25 _____ 15s 607ms/step - accuracy: 0.9430 - loss: 0.1613 - val_accuracy: 0.9245 - val_loss: 0.2478
Epoch 76/100
25/25 _____ 18s 718ms/step - accuracy: 0.9363 - loss: 0.1952 - val_accuracy: 0.9336 - val_loss: 0.2234
Epoch 77/100
25/25 _____ 17s 670ms/step - accuracy: 0.9471 - loss: 0.1735 - val_accuracy: 0.9570 - val_loss: 0.1801
Epoch 78/100

25/25 _____ 17s 703ms/step - accuracy: 0.9369 - loss: 0.1933 - val_accuracy: 0.9688 - val_loss: 0.0760
Epoch 79/100

25/25 _____ 18s 734ms/step - accuracy: 0.9321 - loss: 0.1955 - val_accuracy: 0.9492 - val_loss: 0.1476
Epoch 80/100

25/25 _____ 17s 706ms/step - accuracy: 0.9411 - loss: 0.1811 - val_accuracy: 0.9154 - val_loss: 0.2404
Epoch 81/100

25/25 _____ 18s 719ms/step - accuracy: 0.9276 - loss: 0.1974 - val_accuracy: 0.9414 - val_loss: 0.1991
Epoch 82/100

25/25 _____ 18s 742ms/step - accuracy: 0.9488 - loss: 0.1503 - val_accuracy: 0.9492 - val_loss: 0.1502
Epoch 83/100

25/25 _____ 18s 716ms/step - accuracy: 0.9544 - loss: 0.1544 - val_accuracy: 0.9609 - val_loss: 0.1680
Epoch 84/100

25/25 _____ 17s 698ms/step - accuracy: 0.9523 - loss: 0.1385 - val_accuracy: 0.8997 - val_loss: 0.4473
Epoch 85/100

25/25 _____ 17s 670ms/step - accuracy: 0.9198 - loss: 0.2810 - val_accuracy: 0.9583 - val_loss: 0.1869
Epoch 86/100

25/25 _____ 5s 200ms/step - accuracy: 0.9584 - loss: 0.1381 - val_accuracy: 0.9557 - val_loss: 0.1537
Epoch 87/100

25/25 _____ 28s 894ms/step - accuracy: 0.9450 - loss: 0.1475 - val_accuracy: 0.9466 - val_loss: 0.2282
Epoch 88/100

25/25 _____ 20s 820ms/step - accuracy: 0.9514 - loss: 0.1357 - val_accuracy: 0.9714 - val_loss: 0.0851
Epoch 89/100

25/25 _____ 18s 715ms/step - accuracy: 0.9600 - loss: 0.1269 - val_accuracy: 0.9557 - val_loss: 0.1820
Epoch 90/100

25/25 _____ 19s 774ms/step - accuracy: 0.9453 - loss: 0.1844 - val_accuracy: 0.9663 - val_loss: 0.1307
Epoch 91/100

25/25 _____ 24s 971ms/step - accuracy: 0.9626 - loss: 0.1054 - val_accuracy: 0.9622 - val_loss: 0.1649
Epoch 92/100

25/25 _____ 22s 892ms/step - accuracy: 0.9668 - loss: 0.0992 - val_accuracy: 0.9336 - val_loss: 0.2061
Epoch 93/100


25/25 _____ 20s 820ms/step - accuracy: 0.9382 - loss: 0.2057 - val_accuracy: 0.9466 - val_loss: 0.1889
Epoch 94/100

25/25 _____ 18s 747ms/step - accuracy: 0.9496 - loss:


```
0.1690 - val_accuracy: 0.9688 - val_loss: 0.1536
Epoch 95/100
25/25 _____ 18s 757ms/step - accuracy: 0.9566 - loss:
0.1446 - val_accuracy: 0.9518 - val_loss: 0.2025
Epoch 96/100
25/25 _____ 20s 813ms/step - accuracy: 0.9418 - loss:
0.1785 - val_accuracy: 0.9622 - val_loss: 0.1621
Epoch 97/100
25/25 _____ 17s 703ms/step - accuracy: 0.9659 - loss:
0.0995 - val_accuracy: 0.9544 - val_loss: 0.1786
Epoch 98/100
25/25 _____ 20s 831ms/step - accuracy: 0.9260 - loss:
0.2235 - val_accuracy: 0.9674 - val_loss: 0.0970
Epoch 99/100
25/25 _____ 17s 674ms/step - accuracy: 0.9673 - loss:
0.1011 - val_accuracy: 0.9766 - val_loss: 0.0771
Epoch 100/100
25/25 _____ 17s 672ms/step - accuracy: 0.9515 - loss:
0.1677 - val_accuracy: 0.9714 - val_loss: 0.1154
```

Evaluation

```
plt.figure()
plt.plot(hist.history["loss"],label = "Train Loss", color = "red")
plt.plot(hist.history["val_loss"],label = "Validation Loss", color =
"blue", linestyle="dashed",markeredgecolor = "purple", markeredgewidth
= 2)
plt.title("Model Loss", size = 13)
plt.legend()
plt.show()
```



```
plt.figure()
plt.plot(hist.history["accuracy"],label = "Train Accuracy", color =
"red")
plt.plot(hist.history["val_accuracy"],label = "Validation Accuracy",
color = "blue", linestyle="dashed",markeredgecolor = "purple",
markeredgewidth = 2)
plt.title("Model Accuracy", size = 13)
plt.legend()
plt.show()
```



```
loss, accuracy = model.evaluate(test_generator, steps=800 //
batch_size)

print("Test Loss:", loss)
print("Test Accuracy:", accuracy)
```

12/12 ————— 3s 250ms/step - accuracy: 0.9542 - loss: 0.1452
Test Loss: 0.14572185277938843
Test Accuracy: 0.96484375

Testing

```
def load_and_preprocess_image(filename):
    np_image = Image.open(filename)
    np_image = np_image.resize((100, 100))
    np_image = np.array(np_image).astype('float32') / 255
    np_image = np.expand_dims(np_image, axis=0)
    return np_image

image = load_and_preprocess_image(test_path + "/Mango/0_100.jpg")
```



```
plt.imshow(np.squeeze(image))
plt.axis("off")
plt.show()
```



```
prediction = model.predict(image)
predicted_class = np.argmax(prediction, axis=-1)
print("Predicted class index:", predicted_class)

1/1 ————— 0s 29ms/step
Predicted class index: [64]

class_indices = test_generator.class_indices
predicted_class_label = [label for label, index in
class_indices.items() if index == 64]
print("Predicted class label:", predicted_class_label)

Predicted class label: ['Mango']
```

```
img_path =
'fruits-360_dataset/fruits-360/test-multiple_fruits/Bananas(lady_finge
r)1.jpg'
preprocessed_image = load_and_preprocess_image(img_path)
```

```
plt.imshow(np.squeeze(preprocessed_image), interpolation='nearest')
plt.axis('off')
plt.show()
```



```
prediction = model.predict(preprocessed_image)
predicted_class = np.argmax(prediction, axis=-1)
print("Predicted class index:", predicted_class)

1/1 ————— 0s 23ms/step
Predicted class index: [26]

predicted_class_label = [label for label, index in
class_indices.items() if index == 17]
print("Predicted class label:", predicted_class_label)

Predicted class label: ['Banana Lady Finger']
```

```
img_path =
'fruits-360_dataset/fruits-360/test-multiple_fruits/grape_pear_mandari
ne.jpg'
preprocessed_image = load_and_preprocess_image(img_path)
```

```
plt.imshow(np.squeeze(preprocessed_image), interpolation='nearest')
plt.axis('off')
plt.show()
```



```
prediction = model.predict(preprocessed_image)
predicted_class = np.argmax(prediction, axis=-1)
print("Predicted class index:", predicted_class)

1/1 ————— 0s 23ms/step
Predicted class index: [37]

predicted_class_label = [label for label, index in
class_indices.items() if index == 116]
print("Predicted class label:", predicted_class_label)

Predicted class label: ['Strawberry']
```

Model Evaluation

Based off what we see from the Model evaluation visualizations and training loss and decreases really quickly with the exact opposite effect for the Model Accuracy. By printing out the loss and accuracy we can get a better understanding of the loss and accuracy of the model which the Test Loss: 0.25146737694740295 and the Test Accuracy is 0.9244791865348816. From testing multiple images above we can see a success rate with individual images like the mango but with testing multiple images it gets a little confusing when working with images that have

multiple different fruit (like the grape, pears, and mandarines). However it picks up on multiple of the same fruit fairly well (Banana Lady Finger).

Alt Models

```
input_shape = x.shape

# Batch Normalization Model
def create_bn_model():
    model = Sequential([
        Conv2D(32, (3, 3), input_shape=input_shape),
        Activation('relu'),
        BatchNormalization(),
        MaxPooling2D(),
        Conv2D(64, (3, 3)),
        Activation('relu'),
        BatchNormalization(),
        MaxPooling2D(),
        Conv2D(128, (3, 3)),
        Activation('relu'),
        BatchNormalization(),
        MaxPooling2D(),
        Flatten(),
        Dense(1024, activation='relu'),
        Dropout(0.5),
        Dense(number_of_class, activation='softmax')
    ])
    model.compile(optimizer='rmsprop',
        loss='categorical_crossentropy', metrics=['accuracy'])
    return model

# Adjusted Learning Rate and Additional Dropout
def create_lr_dropout_model():
    model = Sequential([
        Conv2D(32, (3, 3), input_shape=input_shape),
        Activation('relu'),
        MaxPooling2D(),
        Dropout(0.3),
        Conv2D(64, (3, 3)),
        Activation('relu'),
        MaxPooling2D(),
        Dropout(0.3),
        Conv2D(128, (3, 3)),
        Activation('relu'),
        MaxPooling2D(),
        Flatten(),
        Dense(1024, activation='relu'),
```

```

        Dropout(0.5),
        Dense(number_of_class, activation='softmax')
    ])
    optimizer = Adam(learning_rate=0.0005)
    model.compile(optimizer=optimizer,
loss='categorical_crossentropy', metrics=['accuracy'])
    return model

bn_model = create_bn_model()
lr_dropout_model = create_lr_dropout_model()

print("\nBatch Normalization Model Summary:")
bn_model.summary()
print("\nLearning Rate and Dropout Model Summary:")
lr_dropout_model.summary()

```

Batch Normalization Model Summary:

Model: "sequential_2"

Layer (type) Param #	Output Shape	
conv2d_6 (Conv2D) 896	(None, 98, 98, 32)	
activation_10 (Activation) 0	(None, 98, 98, 32)	
batch_normalization 128 (BatchNormalization)	(None, 98, 98, 32)	
max_pooling2d_6 (MaxPooling2D) 0	(None, 49, 49, 32)	
conv2d_7 (Conv2D) 18,496	(None, 47, 47, 64)	
activation_11 (Activation)	(None, 47, 47, 64)	

0				
		batch_normalization_1	(None, 47, 47, 64)	
256		(BatchNormalization)		
		max_pooling2d_7 (MaxPooling2D)	(None, 23, 23, 64)	
0				
		conv2d_8 (Conv2D)	(None, 21, 21, 128)	
73,856				
		activation_12 (Activation)	(None, 21, 21, 128)	
0				
		batch_normalization_2	(None, 21, 21, 128)	
512		(BatchNormalization)		
		max_pooling2d_8 (MaxPooling2D)	(None, 10, 10, 128)	
0				
		flatten_2 (Flatten)	(None, 12800)	
0				
		dense_4 (Dense)	(None, 1024)	
13,108,224				
		dropout_2 (Dropout)	(None, 1024)	
0				
		dense_5 (Dense)	(None, 131)	
134,275				

Total params: 13,336,643 (50.88 MB)

Trainable params: 13,336,195 (50.87 MB)

Non-trainable params: 448 (1.75 KB)

Learning Rate and Dropout Model Summary:

Model: "sequential_3"

Layer (type) Param #	Output Shape
conv2d_9 (Conv2D) 896	(None, 98, 98, 32)
activation_13 (Activation) 0	(None, 98, 98, 32)
max_pooling2d_9 (MaxPooling2D) 0	(None, 49, 49, 32)
dropout_3 (Dropout) 0	(None, 49, 49, 32)
conv2d_10 (Conv2D) 18,496	(None, 47, 47, 64)
activation_14 (Activation) 0	(None, 47, 47, 64)
max_pooling2d_10 (MaxPooling2D) 0	(None, 23, 23, 64)
dropout_4 (Dropout) 0	(None, 23, 23, 64)
conv2d_11 (Conv2D) 73,856	(None, 21, 21, 128)

0	activation_15 (Activation)	(None, 21, 21, 128)	
0	max_pooling2d_11 (MaxPooling2D)	(None, 10, 10, 128)	
0	flatten_3 (Flatten)	(None, 12800)	
13,108,224	dense_6 (Dense)	(None, 1024)	
0	dropout_5 (Dropout)	(None, 1024)	
134,275	dense_7 (Dense)	(None, 131)	

Total params: 13,335,747 (50.87 MB)

Trainable params: 13,335,747 (50.87 MB)

Non-trainable params: 0 (0.00 B)

```

checkpoint_bn = ModelCheckpoint('best_model_bn.keras',
monitor='val_accuracy', mode='max', save_best_only=True)
early_stop_bn = EarlyStopping(monitor='val_loss', patience=10,
restore_best_weights=True)

checkpoint_lr_dropout = ModelCheckpoint('best_model_lr_dropout.keras',
monitor='val_accuracy', mode='max', save_best_only=True)
early_stop_lr_dropout = EarlyStopping(monitor='val_loss', patience=10,
restore_best_weights=True)

# Train the Batch Normalization Model
history_bn = bn_model.fit(
    x=train_generator,
    steps_per_epoch=1600 // batch_size,
    epochs=epochs,
    validation_data=test_generator,

```



```

validation_steps=800 // batch_size,
callbacks=[checkpoint_bn, early_stop_bn]
)

Epoch 1/100
25/25 _____ 29s 1s/step - accuracy: 0.0780 - loss:
7.2950 - val_accuracy: 0.0065 - val_loss: 9.3268
Epoch 2/100
25/25 _____ 25s 1s/step - accuracy: 0.1544 - loss:
5.2696 - val_accuracy: 0.0143 - val_loss: 9.5143
Epoch 3/100
25/25 _____ 25s 1s/step - accuracy: 0.2160 - loss:
4.2411 - val_accuracy: 0.0104 - val_loss: 10.5254
Epoch 4/100
25/25 _____ 25s 1s/step - accuracy: 0.2858 - loss:
3.7232 - val_accuracy: 0.0052 - val_loss: 13.9616
Epoch 5/100
25/25 _____ 27s 1s/step - accuracy: 0.3263 - loss:
3.2050 - val_accuracy: 0.0065 - val_loss: 18.7386
Epoch 6/100
25/25 _____ 27s 1s/step - accuracy: 0.3549 - loss:
3.1484 - val_accuracy: 0.0065 - val_loss: 16.0833
Epoch 7/100
25/25 _____ 22s 906ms/step - accuracy: 0.4018 - loss:
2.7747 - val_accuracy: 0.0117 - val_loss: 17.9076
Epoch 8/100
25/25 _____ 24s 960ms/step - accuracy: 0.4551 - loss:
2.5361 - val_accuracy: 0.0391 - val_loss: 17.1191
Epoch 9/100
25/25 _____ 23s 949ms/step - accuracy: 0.5026 - loss:
2.1727 - val_accuracy: 0.0352 - val_loss: 21.3140
Epoch 10/100
25/25 _____ 23s 951ms/step - accuracy: 0.4857 - loss:
2.4426 - val_accuracy: 0.0234 - val_loss: 18.3315
Epoch 11/100
25/25 _____ 26s 1s/step - accuracy: 0.5790 - loss:
1.7099 - val_accuracy: 0.0508 - val_loss: 14.8973

fig, axs = plt.subplots(nrows=1, ncols=2, figsize=(16, 6))

axs[0].plot(history_bn.history['accuracy'], label='Train Accuracy',
color='red')
axs[0].plot(history_bn.history['val_accuracy'], label='Validation
Accuracy', color='blue', linestyle='dashed', markeredgewidth=2,
markedgewidth=2)
axs[0].set_title('Model Accuracy', size=13)
axs[0].set_xlabel('Epoch')
axs[0].set_ylabel('Accuracy')
axs[0].legend()


```

```

axs[1].plot(history_bn.history['loss'], label='Train Loss',
color='red')
axs[1].plot(history_bn.history['val_loss'], label='Validation Loss',
color='blue', linestyle='dashed', markeredgewidth=2,
markeredgecolor='purple',
markeredgewidth=2)
axs[1].set_title('Model Loss', size=13)
axs[1].set_xlabel('Epoch')
axs[1].set_ylabel('Loss')
axs[1].legend()

# Display the plots
plt.tight_layout()
plt.show()

```



```

final_train_accuracy = history_bn.history['accuracy'][-1]
final_val_accuracy = history_bn.history['val_accuracy'][-1]
final_train_loss = history_bn.history['loss'][-1]
final_val_loss = history_bn.history['val_loss'][-1]

print(f"Final Training Accuracy: {final_train_accuracy:.4f}")
print(f"Final Validation Accuracy: {final_val_accuracy:.4f}")
print(f"Final Training Loss: {final_train_loss:.4f}")
print(f"Final Validation Loss: {final_val_loss:.4f}")

bn_loss, bn_accuracy = bn_model.evaluate(test_generator, steps=800 //
batch_size)
print("Batch Normalization Model - Test Loss:", bn_loss)
print("Batch Normalization Model - Test Accuracy:", bn_accuracy)

Final Training Accuracy: 0.5638
Final Validation Accuracy: 0.0508
Final Training Loss: 1.8425
Final Validation Loss: 14.8973

```

12/12 ————— 3s 235ms/step - accuracy: 0.0069 - loss: 9.1953

Batch Normalization Model - Test Loss: 9.354384422302246

Batch Normalization Model - Test Accuracy: 0.0065104165114462376

Train the Learning Rate and Dropout Model

```
history_lr_dropout = lr_dropout_model.fit(  
    x=train_generator,  
    steps_per_epoch=1600 // batch_size,  
    epochs=epochs,  
    validation_data=test_generator,  
    validation_steps=800 // batch_size,  
    callbacks=[checkpoint_lr_dropout, early_stop_lr_dropout]  
)
```

Epoch 1/100

25/25 ————— 21s 729ms/step - accuracy: 0.4360 - loss: 1.8765 - val_accuracy: 0.6979 - val_loss: 1.4491

Epoch 2/100

25/25 ————— 16s 657ms/step - accuracy: 0.4629 - loss: 1.7675 - val_accuracy: 0.6576 - val_loss: 1.3786

Epoch 3/100

25/25 ————— 15s 633ms/step - accuracy: 0.4831 - loss: 1.6620 - val_accuracy: 0.7109 - val_loss: 1.1578

Epoch 4/100

25/25 ————— 15s 597ms/step - accuracy: 0.5676 - loss: 1.4479 - val_accuracy: 0.7474 - val_loss: 1.1082

Epoch 5/100

25/25 ————— 15s 598ms/step - accuracy: 0.5992 - loss: 1.2889 - val_accuracy: 0.7448 - val_loss: 0.9857

Epoch 6/100

25/25 ————— 15s 627ms/step - accuracy: 0.6163 - loss: 1.2194 - val_accuracy: 0.7578 - val_loss: 1.0530

Epoch 7/100

25/25 ————— 15s 621ms/step - accuracy: 0.6190 - loss: 1.1668 - val_accuracy: 0.7617 - val_loss: 0.8701

Epoch 8/100

25/25 ————— 16s 633ms/step - accuracy: 0.6818 - loss: 1.0526 - val_accuracy: 0.7604 - val_loss: 0.8290

Epoch 9/100

25/25 ————— 16s 662ms/step - accuracy: 0.6808 - loss: 0.9771 - val_accuracy: 0.8086 - val_loss: 0.7572

Epoch 10/100

25/25 ————— 16s 644ms/step - accuracy: 0.7112 - loss: 0.9261 - val_accuracy: 0.8034 - val_loss: 0.7234

Epoch 11/100

25/25 ————— 16s 669ms/step - accuracy: 0.7375 - loss: 0.8374 - val_accuracy: 0.8333 - val_loss: 0.6933

Epoch 12/100

25/25 ————— 17s 678ms/step - accuracy: 0.7474 - loss:

0.7785 - val_accuracy: 0.8646 - val_loss: 0.5132
Epoch 13/100
25/25 _____ 16s 643ms/step - accuracy: 0.7643 - loss:
0.7036 - val_accuracy: 0.8581 - val_loss: 0.5351
Epoch 14/100
25/25 _____ 16s 656ms/step - accuracy: 0.7733 - loss:
0.6941 - val_accuracy: 0.8372 - val_loss: 0.6205
Epoch 15/100
25/25 _____ 17s 676ms/step - accuracy: 0.7628 - loss:
0.7335 - val_accuracy: 0.8281 - val_loss: 0.5734
Epoch 16/100
25/25 _____ 17s 708ms/step - accuracy: 0.7795 - loss:
0.6694 - val_accuracy: 0.8685 - val_loss: 0.4946
Epoch 17/100
25/25 _____ 16s 643ms/step - accuracy: 0.7715 - loss:
0.6904 - val_accuracy: 0.8594 - val_loss: 0.5224
Epoch 18/100
25/25 _____ 16s 648ms/step - accuracy: 0.7736 - loss:
0.6570 - val_accuracy: 0.8802 - val_loss: 0.4269
Epoch 19/100
25/25 _____ 19s 758ms/step - accuracy: 0.8234 - loss:
0.5267 - val_accuracy: 0.8945 - val_loss: 0.3869
Epoch 20/100
25/25 _____ 18s 728ms/step - accuracy: 0.8214 - loss:
0.5637 - val_accuracy: 0.8685 - val_loss: 0.4203
Epoch 21/100
25/25 _____ 19s 753ms/step - accuracy: 0.8348 - loss:
0.4779 - val_accuracy: 0.8750 - val_loss: 0.4527
Epoch 22/100
25/25 _____ 19s 754ms/step - accuracy: 0.8338 - loss:
0.5117 - val_accuracy: 0.8919 - val_loss: 0.3967
Epoch 23/100
25/25 _____ 20s 823ms/step - accuracy: 0.8549 - loss:
0.4923 - val_accuracy: 0.9102 - val_loss: 0.3177
Epoch 24/100
25/25 _____ 19s 756ms/step - accuracy: 0.8487 - loss:
0.4899 - val_accuracy: 0.8516 - val_loss: 0.5551
Epoch 25/100
25/25 _____ 20s 795ms/step - accuracy: 0.8282 - loss:
0.5090 - val_accuracy: 0.8789 - val_loss: 0.4061
Epoch 26/100
25/25 _____ 19s 790ms/step - accuracy: 0.8456 - loss:
0.4416 - val_accuracy: 0.8971 - val_loss: 0.3299
Epoch 27/100
25/25 _____ 25s 1s/step - accuracy: 0.8825 - loss:
0.3597 - val_accuracy: 0.9128 - val_loss: 0.3073
Epoch 28/100
25/25 _____ 22s 888ms/step - accuracy: 0.8598 - loss:
0.3801 - val_accuracy: 0.9297 - val_loss: 0.2569

```

Epoch 29/100
25/25 _____ 16s 668ms/step - accuracy: 0.8766 - loss:
0.3679 - val_accuracy: 0.9115 - val_loss: 0.3112
Epoch 30/100
25/25 _____ 17s 670ms/step - accuracy: 0.8893 - loss:
0.3698 - val_accuracy: 0.8942 - val_loss: 0.3451
Epoch 31/100
25/25 _____ 19s 787ms/step - accuracy: 0.8572 - loss:
0.4376 - val_accuracy: 0.9010 - val_loss: 0.3048
Epoch 32/100
25/25 _____ 18s 754ms/step - accuracy: 0.8837 - loss:
0.3672 - val_accuracy: 0.9102 - val_loss: 0.3315
Epoch 33/100
25/25 _____ 19s 781ms/step - accuracy: 0.8835 - loss:
0.3297 - val_accuracy: 0.8841 - val_loss: 0.3667
Epoch 34/100
25/25 _____ 19s 775ms/step - accuracy: 0.8847 - loss:
0.3652 - val_accuracy: 0.9219 - val_loss: 0.2896
Epoch 35/100
25/25 _____ 22s 912ms/step - accuracy: 0.8887 - loss:
0.3383 - val_accuracy: 0.9180 - val_loss: 0.3337
Epoch 36/100
25/25 _____ 16s 653ms/step - accuracy: 0.8852 - loss:
0.3400 - val_accuracy: 0.9128 - val_loss: 0.3055
Epoch 37/100
25/25 _____ 15s 599ms/step - accuracy: 0.8896 - loss:
0.3411 - val_accuracy: 0.9154 - val_loss: 0.2632
Epoch 38/100
25/25 _____ 15s 620ms/step - accuracy: 0.8837 - loss:
0.3421 - val_accuracy: 0.9180 - val_loss: 0.2806

```

```
fig, axs = plt.subplots(nrows=1, ncols=2, figsize=(16, 6))
```

```
axs[0].plot(history_lr_dropout.history['accuracy'], label='Train
Accuracy', color='red')
```

```
axs[0].plot(history_lr_dropout.history['val_accuracy'],
label='Validation Accuracy', color='blue', linestyle='dashed',
markeredgecolor='purple', markeredgewidth=2)
```

```
axs[0].set_title('Model Accuracy', size=13)
```

```
axs[0].set_xlabel('Epoch')
```

```
axs[0].set_ylabel('Accuracy')
```

```
axs[0].legend()
```

```
axs[1].plot(history_lr_dropout.history['loss'], label='Train Loss',
color='red')
```

```
axs[1].plot(history_lr_dropout.history['val_loss'], label='Validation
Loss', color='blue', linestyle='dashed', markeredgecolor='purple',
markeredgewidth=2)
```

```
axs[1].set_title('Model Loss', size=13)
```

```
axs[1].set_xlabel('Epoch')
```

```


axs[1].set_ylabel('Loss')
axs[1].legend()

```

```

plt.tight_layout()
plt.show()

```



```

final_train_accuracy = history_lr_dropout.history['accuracy'][-1]
final_val_accuracy = history_lr_dropout.history['val_accuracy'][-1]
final_train_loss = history_lr_dropout.history['loss'][-1]
final_val_loss = history_lr_dropout.history['val_loss'][-1]

```

```

print(f"Final Training Accuracy: {final_train_accuracy:.4f}")
print(f"Final Validation Accuracy: {final_val_accuracy:.4f}")
print(f"Final Training Loss: {final_train_loss:.4f}")
print(f"Final Validation Loss: {final_val_loss:.4f}")

```

```

lr_dropout_loss, lr_dropout_accuracy =
lr_dropout_model.evaluate(test_generator, steps=800 // batch_size)
print("Learning Rate and Dropout Model - Test Loss:", lr_dropout_loss)
print("Learning Rate and Dropout Model - Test Accuracy:",
lr_dropout_accuracy)

```

```

Final Training Accuracy: 0.8888
Final Validation Accuracy: 0.9180
Final Training Loss: 0.3339
Final Validation Loss: 0.2806
12/12 2s 137ms/step - accuracy: 0.9161 - loss:
0.3143
Learning Rate and Dropout Model - Test Loss: 0.31704944372177124
Learning Rate and Dropout Model - Test Accuracy: 0.91015625

```

Testing

```

def load_and_preprocess_image(image_path):
    img = load_img(image_path, target_size=(100, 100))
    img_array = img_to_array(img) / 255.0
    img_array = np.expand_dims(img_array, axis=0)
    return img_array

def get_class_label(class_index, class_indices):
    labels = dict((v,k) for k,v in class_indices.items())
    return labels[class_index]

img_path =
'fruits-360_dataset/fruits-360/test-multiple_fruits/mangos1.jpg'
preprocessed_image = load_and_preprocess_image(img_path)
plt.imshow(np.squeeze(preprocessed_image))
plt.axis('off')
plt.show()

```



```

prediction = lr_dropout_model.predict(preprocessed_image)
predicted_class_index = np.argmax(prediction, axis=-1)[0]
print("Predicted class index:", predicted_class_index)

1/1 _____ 0s 78ms/step
Predicted class index: 14

predicted_class_label = get_class_label(predicted_class_index,
train_generator.class_indices)
print("Predicted class label:", predicted_class_label)

```

Predicted class label: Avocado

```
img_path =  
'fruits-360_dataset/fruits-360/test-multiple_fruits/apple.jpg'  
preprocessed_image = load_and_preprocess_image(img_path)  
plt.imshow(np.squeeze(preprocessed_image))  
plt.axis('off')  
plt.show()
```



```
prediction = lr_dropout_model.predict(preprocessed_image)  
predicted_class_index = np.argmax(prediction, axis=-1)[0]  
print("Predicted class index:", predicted_class_index)  
  
1/1 _____ 0s 22ms/step  
Predicted class index: 12  
  
predicted_class_label = get_class_label(predicted_class_index,  
train_generator.class_indices)  
print("Predicted class label:", predicted_class_label)  
  
Predicted class label: Apple Red Yellow 2
```

```
image = test_path + "/Beetroot/23_100.jpg"
preprocessed_image = load_and_preprocess_image(image)
plt.imshow(np.squeeze(preprocessed_image))
plt.axis('off')
plt.show()
```



```
prediction = lr_dropout_model.predict(preprocessed_image)
predicted_class_index = np.argmax(prediction, axis=-1)[0]
print("Predicted class index:", predicted_class_index)

1/1 _____ 0s 26ms/step
Predicted class index: 19

predicted_class_label = get_class_label(predicted_class_index,
train_generator.class_indices)
print("Predicted class label:", predicted_class_label)

Predicted class label: Beetroot
```

Own test pics (for fun)

```
image= "Adam test imgs/ban.png"
preprocessed_image = load_and_preprocess_image(image)
```

```
plt.imshow(np.squeeze(preprocessed_image))
plt.axis('off')
plt.show()
```



```
prediction = lr_dropout_model.predict(preprocessed_image)
predicted_class_index = np.argmax(prediction, axis=-1)[0]
print("Predicted class index:", predicted_class_index)

1/1 _____ 0s 30ms/step
Predicted class index: 17

predicted_class_label = get_class_label(predicted_class_index,
train_generator.class_indices)
print("Predicted class label:", predicted_class_label)

Predicted class label: Banana Lady Finger

image = "Adam test imgs/shrimp.png"

preprocessed_image = load_and_preprocess_image(image)
plt.imshow(np.squeeze(preprocessed_image))
plt.axis('off')
plt.show()
```



```
prediction = lr_dropout_model.predict(preprocessed_image)
predicted_class_index = np.argmax(prediction, axis=-1)[0]
print("Predicted class index:", predicted_class_index)
```

```
1/1 ————— 0s 24ms/step
Predicted class index: 36
```

```
predicted_class_label = get_class_label(predicted_class_index,
train_generator.class_indices)
print("Predicted class label:", predicted_class_label)
```

```
Predicted class label: Corn Husk
```

Maybe its because shrimp isnt in the training set `_(\ツ)/~`

```
def create_improved_model(input_shape, number_of_classes):
    model = Sequential([
        Conv2D(16, (2, 2), padding='same', input_shape=input_shape),
        Activation('relu'),
        MaxPooling2D(pool_size=2),
        Conv2D(32, (2, 2), activation='relu', padding='same'),
        MaxPooling2D(pool_size=2),
        Conv2D(64, (2, 2), activation='relu', padding='same'),
        MaxPooling2D(pool_size=2),
        Conv2D(128, (2, 2), activation='relu', padding='same'),
```

```

        MaxPooling2D(pool_size=2),
        Dropout(0.3),
        Flatten(),
        Dense(150, activation='relu'),
        Dropout(0.4),
        Dense(number_of_classes, activation='softmax')
    ])
    return model

```

```

model = create_improved_model(input_shape=(100, 100, 3),
number_of_classes=number_of_class)
model.compile(optimizer='rmsprop', loss='categorical_crossentropy',
metrics=['accuracy'])

```

```

print("\nBatch Normalization Model Summary:")
model.summary()

```

Batch Normalization Model Summary:

Model: "sequential_4"

Layer (type) Param #	Output Shape	
conv2d_12 (Conv2D) 208	(None, 100, 100, 16)	
activation_16 (Activation) 0	(None, 100, 100, 16)	
max_pooling2d_12 (MaxPooling2D) 0	(None, 50, 50, 16)	
conv2d_13 (Conv2D) 2,080	(None, 50, 50, 32)	
max_pooling2d_13 (MaxPooling2D) 0	(None, 25, 25, 32)	
conv2d_14 (Conv2D) 8,256	(None, 25, 25, 64)	

0	max_pooling2d_14 (MaxPooling2D)	(None, 12, 12, 64)
32,896	conv2d_15 (Conv2D)	(None, 12, 12, 128)
0	max_pooling2d_15 (MaxPooling2D)	(None, 6, 6, 128)
0	dropout_6 (Dropout)	(None, 6, 6, 128)
0	flatten_4 (Flatten)	(None, 4608)
691,350	dense_8 (Dense)	(None, 150)
0	dropout_7 (Dropout)	(None, 150)
19,781	dense_9 (Dense)	(None, 131)

Total params: 1,509,144 (5.76 MB)

Trainable params: 754,571 (2.88 MB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 754,573 (2.88 MB)

```
train_datagen = ImageDataGenerator(
    rescale=1./255,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True
)
```

```

test_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(
    directory=train_path,
    target_size=(100, 100), # Ensure the target size is the same as
    `input_shape`
    batch_size=32,
    class_mode='categorical'
)

test_generator = test_datagen.flow_from_directory(
    directory=test_path,
    target_size=(100, 100),
    batch_size=32,
    class_mode='categorical'
)

Found 67692 images belonging to 131 classes.
Found 22688 images belonging to 131 classes.

history = model.fit(
    train_generator,
    steps_per_epoch=1600 // 32,
    epochs=30,
    validation_data=test_generator,
    validation_steps=800 // 32
)

Epoch 1/30
50/50 _____ 22s 405ms/step - accuracy: 0.0100 - loss:
4.8697 - val_accuracy: 0.0188 - val_loss: 4.8105
Epoch 2/30
50/50 _____ 18s 363ms/step - accuracy: 0.0201 - loss:
4.7297 - val_accuracy: 0.0463 - val_loss: 4.3688
Epoch 3/30
50/50 _____ 18s 361ms/step - accuracy: 0.0442 - loss:
4.4008 - val_accuracy: 0.1187 - val_loss: 3.7958
Epoch 4/30
50/50 _____ 16s 333ms/step - accuracy: 0.0830 - loss:
3.9222 - val_accuracy: 0.1750 - val_loss: 3.5075
Epoch 5/30
50/50 _____ 13s 256ms/step - accuracy: 0.1137 - loss:
3.6508 - val_accuracy: 0.2788 - val_loss: 3.0553
Epoch 6/30
50/50 _____ 13s 260ms/step - accuracy: 0.1511 - loss:
3.3337 - val_accuracy: 0.3438 - val_loss: 2.6870
Epoch 7/30
50/50 _____ 14s 275ms/step - accuracy: 0.1811 - loss:
3.0400 - val_accuracy: 0.4162 - val_loss: 2.3219

```

Epoch 8/30
50/50 _____ 13s 265ms/step - accuracy: 0.2497 - loss: 2.7689 - val_accuracy: 0.2450 - val_loss: 2.6393

Epoch 9/30
50/50 _____ 13s 265ms/step - accuracy: 0.2904 - loss: 2.5951 - val_accuracy: 0.4250 - val_loss: 2.0041

Epoch 10/30
50/50 _____ 13s 255ms/step - accuracy: 0.3485 - loss: 2.3020 - val_accuracy: 0.4563 - val_loss: 1.8233

Epoch 11/30
50/50 _____ 14s 274ms/step - accuracy: 0.3686 - loss: 2.1273 - val_accuracy: 0.6700 - val_loss: 1.3654

Epoch 12/30
50/50 _____ 13s 266ms/step - accuracy: 0.4079 - loss: 2.0002 - val_accuracy: 0.6087 - val_loss: 1.4417

Epoch 13/30
50/50 _____ 13s 258ms/step - accuracy: 0.4121 - loss: 1.9040 - val_accuracy: 0.6200 - val_loss: 1.2504

Epoch 14/30
50/50 _____ 12s 246ms/step - accuracy: 0.4691 - loss: 1.7533 - val_accuracy: 0.7075 - val_loss: 1.0530

Epoch 15/30
50/50 _____ 11s 232ms/step - accuracy: 0.5144 - loss: 1.6541 - val_accuracy: 0.7262 - val_loss: 1.0454

Epoch 16/30
50/50 _____ 13s 261ms/step - accuracy: 0.5234 - loss: 1.4869 - val_accuracy: 0.6963 - val_loss: 1.0630

Epoch 17/30
50/50 _____ 11s 230ms/step - accuracy: 0.5367 - loss: 1.4758 - val_accuracy: 0.7525 - val_loss: 0.9088

Epoch 18/30
50/50 _____ 11s 216ms/step - accuracy: 0.5754 - loss: 1.2621 - val_accuracy: 0.7788 - val_loss: 0.8164

Epoch 19/30
50/50 _____ 11s 215ms/step - accuracy: 0.6259 - loss: 1.2302 - val_accuracy: 0.7675 - val_loss: 0.8469

Epoch 20/30
50/50 _____ 12s 242ms/step - accuracy: 0.6363 - loss: 1.0980 - val_accuracy: 0.7962 - val_loss: 0.7364

Epoch 21/30
50/50 _____ 11s 226ms/step - accuracy: 0.6318 - loss: 1.1893 - val_accuracy: 0.6637 - val_loss: 1.0925

Epoch 22/30
50/50 _____ 11s 220ms/step - accuracy: 0.6654 - loss: 1.1253 - val_accuracy: 0.7613 - val_loss: 0.7678

Epoch 23/30
50/50 _____ 11s 220ms/step - accuracy: 0.6778 - loss: 1.0192 - val_accuracy: 0.8325 - val_loss: 0.5632


Epoch 24/30

```
50/50 _____ 10s 201ms/step - accuracy: 0.6976 - loss:
0.9627 - val_accuracy: 0.8213 - val_loss: 0.5988
Epoch 25/30
50/50 _____ 10s 201ms/step - accuracy: 0.7020 - loss:
0.9555 - val_accuracy: 0.8112 - val_loss: 0.6308
Epoch 26/30
50/50 _____ 10s 200ms/step - accuracy: 0.7252 - loss:
0.8156 - val_accuracy: 0.8300 - val_loss: 0.5948
Epoch 27/30
50/50 _____ 11s 231ms/step - accuracy: 0.7383 - loss:
0.8273 - val_accuracy: 0.7887 - val_loss: 0.6722
Epoch 28/30
50/50 _____ 10s 209ms/step - accuracy: 0.7360 - loss:
0.7699 - val_accuracy: 0.8400 - val_loss: 0.5895
Epoch 29/30
50/50 _____ 10s 201ms/step - accuracy: 0.7282 - loss:
0.8008 - val_accuracy: 0.7778 - val_loss: 0.6918
Epoch 30/30
50/50 _____ 12s 247ms/step - accuracy: 0.7691 - loss:
0.6893 - val_accuracy: 0.8525 - val_loss: 0.5375
```

```
test_loss, test_accuracy = model.evaluate(test_generator)
print("Test Accuracy:", test_accuracy)
print("Test Loss:", test_loss)
```

```
709/709 _____ 108s 152ms/step - accuracy: 0.8441 -
loss: 0.5282
Test Accuracy: 0.8421191573143005
Test Loss: 0.5355644226074219
```

```
plt.figure()
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend()
plt.show()
```

Redo Preprocessing

we restarting []

```
from sklearn.datasets import load_files

train_dir = 'fruits-360_dataset/fruits-360/Training/'
test_dir = 'fruits-360_dataset/fruits-360/Test/'

def load_dataset(data_path):
    data_loading = load_files(data_path)
    files_add = np.array(data_loading['filenames'])
    targets_fruits = np.array(data_loading['target'])
    target_labels_fruits = np.array(data_loading['target_names'])
    return files_add, targets_fruits, target_labels_fruits

x_train, y_train, target_labels = load_dataset(train_dir)
x_test, y_test, _ = load_dataset(test_dir)

from tensorflow.keras.utils import to_categorical

no_of_classes = len(np.unique(y_train))
```

```

y_train = to_categorical(y_train, num_classes=no_of_classes)
y_test = to_categorical(y_test, num_classes=no_of_classes)

print(y_train[0])

[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.
0.
0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.
0.
0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.
0.
0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.
0.
0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.
0.
0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]

x_test,x_valid = x_test[7000:],x_test[:7000]
y_test,y_valid = y_test[7000:],y_test[:7000]
print('Valiation X : ',x_valid.shape)
print('Validation y : ',y_valid.shape)
print('Test X : ',x_test.shape)
print('Test y : ',y_test.shape)

Vaildation X : (7000,)
Vaildation y : (7000, 131)
Test X : (15688,)
Test y : (15688, 131)

def convert_image_to_array_form(files):
    images_array=[]
    for file in files:
        images_array.append(img_to_array(load_img(file)))
    return images_array

x_train = np.array(convert_image_to_array_form(x_train))
print('Training set shape : ',x_train.shape)

x_valid = np.array(convert_image_to_array_form(x_valid))
print('Validation set shape : ',x_valid.shape)

x_test = np.array(convert_image_to_array_form(x_test))
print('Test set shape : ',x_test.shape)

print('1st training image shape ',x_train[0].shape)

Training set shape : (67692, 100, 100, 3)
Validation set shape : (7000, 100, 100, 3)
Test set shape : (15688, 100, 100, 3)
1st training image shape (100, 100, 3)

```

```
# Tryna reduce training time
x_train = x_train.astype('float32')/255
x_valid = x_valid.astype('float32')/255
x_test = x_test.astype('float32')/255
```

Now lets try this again

```
def tensorflow_based_model():
    model = Sequential()
    model.add(Conv2D(filters=16, kernel_size=2, input_shape=(100, 100,
3), padding='same'))
    model.add(Activation('relu'))
    model.add(MaxPooling2D(pool_size=2))

    model.add(Conv2D(filters=32, kernel_size=2, activation='relu',
padding='same'))
    model.add(MaxPooling2D(pool_size=2))

    model.add(Conv2D(filters=64, kernel_size=2, activation='relu',
padding='same'))
    model.add(MaxPooling2D(pool_size=2))

    model.add(Conv2D(filters=128, kernel_size=2, activation='relu',
padding='same'))
    model.add(MaxPooling2D(pool_size=2))

    model.add(Dropout(0.3))
    model.add(Flatten())

    model.add(Dense(150))
    model.add(Activation('relu'))
    model.add(Dropout(0.4))

    model.add(Dense(no_of_classes, activation='softmax'))

    return model
```

```
model = tensorflow_based_model()
model.compile(loss='categorical_crossentropy', optimizer='rmsprop',
metrics=['accuracy'])
```

```
c:\Users\adamn\AppData\Local\Programs\Python\Python311\Lib\site-
packages\keras\src\layers\convolutional\base_conv.py:99: UserWarning:
Do not pass an `input_shape`/`input_dim` argument to a layer. When
using Sequential models, prefer using an `Input(shape)` object as the
first layer in the model instead.
```

```
    super().__init__()
```

```
bn_model = create_bn_model()
lr_dropout_model = create_lr_dropout_model()
```

```
print("\nBatch Normalization Model Summary:")
bn_model.summary()
```

Batch Normalization Model Summary:

Model: "sequential_6"

Layer (type)	Output Shape	
Param #		
conv2d_20 (Conv2D)	(None, 98, 98, 32)	
896		
activation_19 (Activation)	(None, 98, 98, 32)	
0		
batch_normalization_3	(None, 98, 98, 32)	
128		
(BatchNormalization)		
max_pooling2d_20 (MaxPooling2D)	(None, 49, 49, 32)	
0		
conv2d_21 (Conv2D)	(None, 47, 47, 64)	
18,496		
activation_20 (Activation)	(None, 47, 47, 64)	
0		
batch_normalization_4	(None, 47, 47, 64)	
256		
(BatchNormalization)		
max_pooling2d_21 (MaxPooling2D)	(None, 23, 23, 64)	
0		

conv2d_22 (Conv2D)	(None, 21, 21, 128)	
73,856		
activation_21 (Activation)	(None, 21, 21, 128)	
0		
batch_normalization_5	(None, 21, 21, 128)	
512		
(BatchNormalization)		
max_pooling2d_22 (MaxPooling2D)	(None, 10, 10, 128)	
0		
flatten_6 (Flatten)	(None, 12800)	
0		
dense_12 (Dense)	(None, 1024)	
13,108,224		
dropout_10 (Dropout)	(None, 1024)	
0		
dense_13 (Dense)	(None, 131)	
134,275		

Total params: 13,336,643 (50.88 MB)

Trainable params: 13,336,195 (50.87 MB)

Non-trainable params: 448 (1.75 KB)

```
history = model.fit(x_train,y_train,
                    batch_size = 32,
                    epochs=30,
                    validation_data=(x_valid, y_valid),
                    verbose=2, shuffle=True)
```

Epoch 1/30

2116/2116 - 112s - 53ms/step - accuracy: 0.7025 - loss: 1.0969 -
val_accuracy: 0.8886 - val_loss: 0.3891

Epoch 2/30
2116/2116 - 100s - 47ms/step - accuracy: 0.9535 - loss: 0.1405 -
val_accuracy: 0.9563 - val_loss: 0.1892

Epoch 3/30
2116/2116 - 93s - 44ms/step - accuracy: 0.9714 - loss: 0.0886 -
val_accuracy: 0.9713 - val_loss: 0.1428

Epoch 4/30
2116/2116 - 82s - 39ms/step - accuracy: 0.9784 - loss: 0.0689 -
val_accuracy: 0.9694 - val_loss: 0.1639

Epoch 5/30
2116/2116 - 96s - 45ms/step - accuracy: 0.9823 - loss: 0.0575 -
val_accuracy: 0.9596 - val_loss: 0.2677

Epoch 6/30
2116/2116 - 90s - 42ms/step - accuracy: 0.9849 - loss: 0.0523 -
val_accuracy: 0.9801 - val_loss: 0.1432

Epoch 7/30
2116/2116 - 93s - 44ms/step - accuracy: 0.9865 - loss: 0.0474 -
val_accuracy: 0.9741 - val_loss: 0.1606

Epoch 8/30
2116/2116 - 91s - 43ms/step - accuracy: 0.9874 - loss: 0.0443 -
val_accuracy: 0.9764 - val_loss: 0.1948

Epoch 9/30
2116/2116 - 88s - 41ms/step - accuracy: 0.9890 - loss: 0.0404 -
val_accuracy: 0.9813 - val_loss: 0.1747

Epoch 10/30
2116/2116 - 91s - 43ms/step - accuracy: 0.9894 - loss: 0.0393 -
val_accuracy: 0.9856 - val_loss: 0.1673

Epoch 11/30
2116/2116 - 88s - 42ms/step - accuracy: 0.9901 - loss: 0.0375 -
val_accuracy: 0.9797 - val_loss: 0.1998

Epoch 12/30
2116/2116 - 106s - 50ms/step - accuracy: 0.9906 - loss: 0.0346 -
val_accuracy: 0.9813 - val_loss: 0.1816

Epoch 13/30
2116/2116 - 130s - 62ms/step - accuracy: 0.9914 - loss: 0.0372 -
val_accuracy: 0.9829 - val_loss: 0.1783

Epoch 14/30
2116/2116 - 105s - 50ms/step - accuracy: 0.9911 - loss: 0.0369 -
val_accuracy: 0.9830 - val_loss: 0.1718

Epoch 15/30
2116/2116 - 120s - 57ms/step - accuracy: 0.9917 - loss: 0.0334 -
val_accuracy: 0.9821 - val_loss: 0.1979

Epoch 16/30
2116/2116 - 115s - 54ms/step - accuracy: 0.9922 - loss: 0.0355 -
val_accuracy: 0.9809 - val_loss: 0.2063

Epoch 17/30
2116/2116 - 115s - 54ms/step - accuracy: 0.9919 - loss: 0.0348 -
val_accuracy: 0.9867 - val_loss: 0.1920

Epoch 18/30

```
2116/2116 - 113s - 54ms/step - accuracy: 0.9925 - loss: 0.0342 -  
val_accuracy: 0.9763 - val_loss: 0.2792  
Epoch 19/30  
2116/2116 - 116s - 55ms/step - accuracy: 0.9922 - loss: 0.0337 -  
val_accuracy: 0.9844 - val_loss: 0.1519  
Epoch 20/30  
2116/2116 - 116s - 55ms/step - accuracy: 0.9927 - loss: 0.0325 -  
val_accuracy: 0.9809 - val_loss: 0.2300  
Epoch 21/30  
2116/2116 - 116s - 55ms/step - accuracy: 0.9928 - loss: 0.0327 -  
val_accuracy: 0.9860 - val_loss: 0.2087  
Epoch 22/30  
2116/2116 - 104s - 49ms/step - accuracy: 0.9932 - loss: 0.0328 -  
val_accuracy: 0.9760 - val_loss: 0.2334  
Epoch 23/30  
2116/2116 - 102s - 48ms/step - accuracy: 0.9938 - loss: 0.0279 -  
val_accuracy: 0.9829 - val_loss: 0.2501  
Epoch 24/30  
2116/2116 - 143s - 67ms/step - accuracy: 0.9935 - loss: 0.0342 -  
val_accuracy: 0.9817 - val_loss: 0.2641  
Epoch 25/30  
2116/2116 - 107s - 51ms/step - accuracy: 0.9932 - loss: 0.0348 -  
val_accuracy: 0.9804 - val_loss: 0.2666  
Epoch 26/30  
2116/2116 - 114s - 54ms/step - accuracy: 0.9931 - loss: 0.0369 -  
val_accuracy: 0.9803 - val_loss: 0.2628  
Epoch 27/30  
2116/2116 - 119s - 56ms/step - accuracy: 0.9927 - loss: 0.0355 -  
val_accuracy: 0.9854 - val_loss: 0.3001  
Epoch 28/30  
2116/2116 - 107s - 51ms/step - accuracy: 0.9942 - loss: 0.0320 -  
val_accuracy: 0.9830 - val_loss: 0.2849  
Epoch 29/30  
2116/2116 - 111s - 52ms/step - accuracy: 0.9941 - loss: 0.0322 -  
val_accuracy: 0.9834 - val_loss: 0.3201  
Epoch 30/30  
2116/2116 - 116s - 55ms/step - accuracy: 0.9935 - loss: 0.0358 -  
val_accuracy: 0.9867 - val_loss: 0.2658
```

```
plt.figure(figsize=(12, 5))
```

```
plt.subplot(1, 2, 1)
```

```
plt.plot(history.history['accuracy'])
```

```
plt.plot(history.history['val_accuracy'])
```

```
plt.title('Model Accuracy')
```

```
plt.ylabel('Accuracy')
```


```
plt.xlabel('Epoch')
```

```
plt.legend(['Train', 'Validation'], loc='upper left')
```

```
plt.subplot(1, 2, 2)
```

```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')

plt.show()
```



```
acc_score = model.evaluate(x_test, y_test)
print("Test Loss:", acc_score[0])
print("Test Accuracy:", acc_score[1])
```

491/491 ————— 10s 19ms/step - accuracy: 0.9876 - loss: 0.2267

Test Loss: 0.24676187336444855

Test Accuracy: 0.9871239066123962

jupyter nbconvert --to html FinalProject.ipynb

Cell In[1], line 1

jupyter nbconvert --to html FinalProject.ipynb

SyntaxError: invalid syntax